

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <math.h>
#include "mpi.h"

typedef struct
{
    int p;
    MPI_Comm comm;
    MPI_Comm row_comm;
    MPI_Comm col_comm;
    int q;
    int my_row;
    int my_col;
    int my_rank;
}
GRID_INFO_T;

main(int argc, char* argv[])
{
    int rank;
    int p;
    float *block_A;
    float *block_B;
    float *mat_C;
    int ma, na;
    int mb, nb;
    int i, j;
    GRID_INFO_T grid;
    MPI_Status status;
    FILE *fp;
    double start1, start2, start3, start4;
    double finish1, finish2, finish3, finish4;

    void Setup_grid(GRID_INFO_T* grid);
    void Read_matrix(char* prompt, float block[], int m, int n, GRID_INFO_T* grid);
    float *parallel_Fox(float block_A[], float block_B[], int ma, int nb, int na,
    int mb, GRID_INFO_T* grid);

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    Setup_grid(&grid);

    /* read dimension of matrix A*/
    if (rank == 0)
    {
        //printf("Enter the dimension m, n of the m x n matrix A:\n");
        scanf("%d %d", &ma, &na);
    }

    /*broadcast the demon of matrix A*/
    MPI_Bcast(&ma, 1, MPI_INT, 0, MPI_COMM_WORLD);
    MPI_Bcast(&na, 1, MPI_INT, 0, MPI_COMM_WORLD);

    block_A = (float*)calloc((na * ma / grid.p), sizeof(float));

    /* read matrix A*/
    start1 = MPI_Wtime();
    Read_matrix("The matrix A", block_A, ma, na, &grid);
    finish1 = MPI_Wtime();

```

```

    /* read dimension of matrix B*/
    if (grid.my_rank == 0)
    {
        //printf("Enter the dimension m, n of the m x n matrix B:\n");
        scanf("%d %d", &mb, &nb);
    }

    /*broadcast the demon of matrix B*/
    MPI_Bcast(&mb, 1, MPI_INT, 0, MPI_COMM_WORLD);
    MPI_Bcast(&nb, 1, MPI_INT, 0, MPI_COMM_WORLD);

    if (na != mb)
    {
        printf("Matrices dimension do not match!\n");
        exit;
    }

    block_B = (float*)calloc(nb * mb / grid.p, sizeof(float));
    mat_C = (float*)malloc(ma * nb * sizeof(float));

    /* read matrix B*/
    start2 = MPI_Wtime();
    Read_matrix("The matrix B", block_B, mb, nb, &grid);
    finish2 = MPI_Wtime();

    /* FOX algorithm*/
    start3 = MPI_Wtime();
    mat_C = parallel_Fox(block_A, block_B, ma, nb, na, mb, &grid);
    finish3 = MPI_Wtime();

    /* save result in a file*/
    if (grid.my_rank == 0)
    {
        fp = fopen("C.txt", "a"); //create C.txt

        start4 = MPI_Wtime();
        fprintf(fp, "%d\n", ma);
        fprintf(fp, "%d\n", nb);
        for (i = 0; i < ma * nb; i++)
        {
            fprintf(fp, "%f\n", mat_C[i]);
        }
        finish4 = MPI_Wtime();

        fclose(fp);

        printf("Number of processes: %d\n", grid.p);
        printf("Time elapsed with I/O: %e\n", finish1 + finish2 + finish3 +
        finish4 - start1 - start2 - start3 - start4);
        printf("Time elapsed without I/O: %e\n", finish3 - start3);
    }

    free(block_A);
    free(block_B);
    free(mat_C);

    MPI_Finalize();
}

/* read matrix*/
void Read_matrix(char* prompt, float block[], int m, int n, GRID_INFO_T* grid)
{
    float *temp;
    int m_bar;

```

```

int n_bar;
int i, j, k, h, x;
int l;
int q, dest;
MPI_Status status;

m_bar = m / grid->q;
n_bar = n / grid->q;

temp = (float*)malloc(m * n / grid->q * sizeof(float));

/* this code use MPI_Scatter is shorter, but it needs to allocate entire
matrix, which is memory cost */
/* if (my_rank ==
0)

{
    for (i = 0; i < q; i++) // row of processes
        for (j = 0; j < m; j++) // number of reading when read the part of matrix
            for each row of porc.
                for (k = 0; k < n_bar; k++)
                {
                    scanf("%f", &temp[m * n / p * (j % q) + j / q * n_bar + k + i * m_bar
* n_bar * q]); // rearrange the index
                    // the size of temp equals to the size of the entire matrix
                }
            MPI_Scatter(temp, m * n / p, MPI_FLOAT, block, m * n / p, MPI_FLOAT, 0,
MPI_COMM_WORLD);
        }

/* this code reads 1/q of entries each time then send them to the corresponding
row of processes*/
/* this code is longer, but needs less memory*/
if (grid->my_rank == 0)
{
    for (i = 0; i < grid->q; i++) // row of processes
    {
        for (j = 0; j < m; j++) // number of reading when read the part of matrix
        for each row of porc.
        {
            for (k = 0; k < n_bar; k++)
            {
                scanf("%f", &temp[m_bar * n_bar * (j % grid->q) + j / grid->q * n_bar +
k]); // rearrange the index
            }
        }
    }

    if (i == 0)
    {
        memcpy(block, temp, m_bar * n_bar * sizeof(float)); // keep for proc. 0
    }

    itself
    processes
    for (l = 1; l < grid->q; l++) // send the rest to the rest 1st row
    {
        dest = l;
        MPI_Send(temp + l * n_bar * m_bar, n_bar * m_bar, MPI_FLOAT, dest, 0,
grid->comm);
    }
}
else
{
    for (l = 0; l < grid->q; l++) // send to other processes row wise

```

```

{
    dest = l + i * grid->q;
    MPI_Send(temp + l * n_bar * m_bar, n_bar * m_bar, MPI_FLOAT, dest, 0,
grid->comm);
}
}
}

else
{
    MPI_Recv(block, n_bar * m_bar, MPI_FLOAT, 0, 0, grid->comm, &status);
}

free(temp);
}

/*local dot operation in the block*/
float *local_dot(float block_A[], float block_B[], int ma, int nb, int na,
GRID_INFO_T* grid)
{
    int i, j, k, h;
    int ma_bar, nb_bar, na_bar;
    float *res;

    res = (float*)calloc(ma * nb / grid->p, sizeof(float));
    ma_bar = ma / grid->q;
    nb_bar = nb / grid->q;
    na_bar = na / grid->q;

    for (j = 0; j < ma_bar; j++) // number of rows in block A
    {
        for (k = 0; k < nb_bar; k++) // number of columns in block B
        {
            for (h = 0; h < na_bar; h++)
            {
                res[j * nb_bar + k] = res[j * nb_bar + k] + block_A[h + j * na_bar] *
block_B[k + h * nb_bar];
            }
        }
    }

    return res;
}

/* circular shift in column communicators */
float column_circular_shift(float block_B[], int mb, int nb, GRID_INFO_T* grid)
{
    int dest;
    int source;
    int tag = 0;
    MPI_Status status;

    source = (grid->q + grid->my_row + 1) % grid->q;
    dest = (grid->q + grid->my_row - 1) % grid->q;

    MPI_Sendrecv_replace(block_B, mb * nb / grid->p, MPI_FLOAT, dest, tag, source,
tag, grid->col_comm, &status);
}

/* broadcast in row communicator */
float row_broadcast(float block_A[], float new_block[], int na, int step,
GRID_INFO_T* grid)
{
    int root;
    int count;

```

```

count = ma * na / grid->p;

// determine which block should be broadcasted in step(0....grid.q-1) in each
row
if (grid->my_rank == grid->my_row * grid->q + (grid->my_row + step) % grid->q)
{
    memcpy(new_block, block_A, count * sizeof(float));
}

root = (grid->my_row + step % grid->q) % grid->q;
MPI_Bcast(new_block, count, MPI_FLOAT, root, grid->row_comm);
}

/* parallel Fox algorithm */
float *parallel_Fox(float block_A[],
                   float block_B[],
                   int ma,
                   int nb,
                   int na,
                   int mb,
                   GRID_INFO_T* grid)
{
    int i, j, k, h, max_step;
    float *res;
    float *new_block;
    float *pieces_res;
    float *mat_C;

    float *local_dot(float block_A[], float block_B[], int ma, int nb, int na,
                    float column_circular_shift(float block_A[], int mb, int nb, GRID_INFO_T* grid);
    float row_broadcast(float block_A[], float new_block[], int na, int ma, int
step, GRID_INFO_T* grid);
    void rearrange_result(float *block_C, float *res, int ma, int nb, GRID_INFO_T*
grid);

    max_step = grid->q;
    res = (float*)calloc(ma * nb / grid->p, sizeof(float));
    pieces_res = (float*)malloc(ma * nb / grid->p * sizeof(float));
    new_block = (float*)calloc(na * ma / grid->p, sizeof(float));
    mat_C = (float*)malloc(ma * nb * sizeof(float));

    for (i = 0; i < max_step; i++)
    {
        row_broadcast(block_A, new_block, na, ma, i, grid);
        pieces_res = local_dot(new_block, block_B, ma, nb, na, grid);
        column_circular_shift(block_B, mb, nb, grid);

        for (j = 0; j < ma * nb / grid->p; j++)
        {
            res[j] = res[j] + pieces_res[j]; // update local result
        }
    }

    rearrange_result(res, mat_C, ma, nb, grid); // arrange entries to a normal order

    free(res);
    free(pieces_res);
    free(new_block);

    return mat_C;
}

void Setup_grid( GRID_INFO_T* grid /* out */)
{

```

```

int old_rank;
int dimensions[2];
int wrap_around[2];
int coordinates[2];
int free_coords[2];

/* Set up Global Grid Information */
MPI_Comm_size(MPI_COMM_WORLD, &(grid->p));
MPI_Comm_rank(MPI_COMM_WORLD, &(grid->rank));

/* We assume p is a perfect square */
grid->q = (int) sqrt((double) grid->p);
dimensions[0] = dimensions[1] = grid->q;

/* We want a circular shift in second dimension. */
/* Don't care about first
wrap_around[0] = wrap_around[1] = 1;
MPI_Cart_create(MPI_COMM_WORLD, 2, dimensions,
                wrap_around, 1, &(grid->comm));
MPI_Comm_rank(grid->comm, &(grid->my_rank));
MPI_Cart_coords(grid->comm, grid->my_rank, 2,
                coordinates);
grid->my_row = coordinates[0];
grid->my_col = coordinates[1];

/* Set up row communicators */
free_coords[0] = 0;
free_coords[1] = 1;
MPI_Cart_sub(grid->comm, free_coords,
              &(grid->row_comm));

/* Set up column communicators */
free_coords[0] = 1;
free_coords[1] = 0;
MPI_Cart_sub(grid->comm, free_coords,
              &(grid->col_comm));
} /* Setup_grid */

void rearrange_result(float *block_C, float *res, int ma, int nb, GRID_INFO_T*
grid)
{
    float *buff;
    int i, j, k;
    int nb_bar, ma_bar;

    ma_bar = ma / grid->q;
    nb_bar = nb / grid->q;
    buff = (float*)malloc(ma_bar * nb * sizeof(float));

    for (i = 0; i < ma_bar; i++) // rearrange the entries in each row of grid
        processes
    {
        MPI_Gather(block_C + i * nb_bar, nb_bar, MPI_FLOAT, buff + i * nb, nb_bar,
MPI_FLOAT, 0, grid->row_comm);
    }

    // gather the ordered entries from first column to the grid processes in proc.0
    MPI_Gather(buff, ma_bar * nb, MPI_FLOAT, res, ma_bar * nb, MPI_FLOAT, 0, grid-
>col_comm);

    free(buff);
}

```